

Matthew Dong

matthew.dong@nyu.edu | linkedin.com/in/matt-j-dong | matt-j-dong.github.io

SUMMARY

Top 10% Master's graduate at **New York University** with **3+** years of **Python** and **PyTorch** experience, specializing in **DevOps** and **CI/CD** automation using **GitHub Actions** to streamline code pipelines. Experienced in cloud computing with **Google Cloud Platform (GCP)**. Proficient in fine-tuning open-source transformer **LLMs** to personalize them and training models to avoid hallucination by using **RAG**. Well-versed in containerization and orchestration using **Docker** and **Kubernetes (K8s)**, as well as **PostgreSQL** data visualization tools such as **Domo** and **Tableau**.

EDUCATION

New York University <i>Master of Science: Computer Science</i>	New York, NY GPA: 3.95/4.00 — Sep 2024 – May 2025
New York University <i>Bachelor of Arts: Computer Science; Bachelor of Arts: Data Science</i>	New York, NY GPA: 3.93/4.00 — Sep 2021 – May 2024

EXPERIENCE

NJStat <i>Artificial Intelligence Intern</i>	Bridgewater, NJ Jan 2025 – Present
--	--

Fast-growing statistical programming startup with demonstrated success in more than 100 clinical trials.

- Updated natural language processing AI using the Transformers library, improved employee monitoring by 20%.
- Leveraged zero-shot and chain-of-thought prompt engineering techniques in lieu of fine-tuning to overcome limited computing resources and drive project success.
- Demonstrated the value of developed applications to management with Figma prototypes, accelerating project approvals and maintaining consistent communication between the team and management.
- Revamped the company timesheet application by designing a robust SQL database schema, processing data, implementing full-stack (front-end and back-end) functionality, and authoring comprehensive documentation for future maintenance.

Colgate-Palmolive <i>Application and Software Development Intern (E-Commerce DevOps)</i>	Piscataway, NJ Jun 2022 – Dec 2024
--	--

Leading e-commerce consumer products company, serving both direct-to-consumer and business-to-business markets.

- Developed Python scripts that utilize cloud functions and API gateways to parse JSON alerts for various applications and send them to Google Chat spaces using webhooks, clearing cluttered email inboxes.
- Automated code updates by implementing robust CI/CD workflows using GitHub Actions while also maximizing reusability—deploying Python scripts to Google Cloud Platform as cloud functions that interface with API gateways.
- Streamlined Splunk panels to convey key information about company servers to developers in an easy to view dashboard.
- Investigated the benefits of using computer vision models such as DALL-E 3 for creating internal company diagrams by implementing a prototype with the OpenAI API.

Sanofi <i>Data Analyst Intern</i>	Bridgewater, NJ Dec 2020 – Feb 2021
---	---

Well-known global pharmaceutical company working to develop new drugs through extensive data analysis.

- Applied PCA, random forests, and K-means clustering to analyze and extract insights from high-dimensional financial and economic data.
- Utilized R to analyze COVID-19 data, applying linear, logistic, and multiple regression techniques to derive insights.

PERSONAL PROJECTS

(More projects available on GitHub: github.com/Matt-J-Dong)

VentureAI (Python: PyTorch)

- Developed an automated travel agent that utilizes fine-tuned open-source LLMs that provides personalized travel recommendations tailored to each user's preferences by integrating the LLM with APIs for dynamic retrieval augmented generation (Dynamic RAG)
- Addressed limitations of state of the art models (hallucination) and implemented innovations such as LoRA and quantization.

Multicore Chess (C++: OpenMP)

- Greatly improved the speed of move calculations, achieving up to 12x speedup in comparison to the sequential version.
- Optimized the evaluation function and parallel alpha-beta pruning algorithms despite limited CPUs and no access to GPUs.

Top-Towns-To-Take-Over-Tech (Spark Scala, Java MapReduce, Shell)

- Built an analytic with a group to find out which American cities are the best for tech careers.
- We cleaned, profiled, and performed ETL on our data before joining all the datasets together, and displayed our results through a Tableau visualization and a Hive warehouse.

SKILLS

Languages: Python, R, Markdown, C++, Spark Scala, MySQL, R, Bash, Java, x86-64 Assembly

Libraries: PyTorch, Transformers, Wandb, NumPy, Pandas, Matplotlib, Seaborn, Scikit-Learn, OpenMP, MPI

Technologies: GitHub, GitHub Actions, Hadoop, Kubernetes (K8s), Docker, GCP, Terraform, Jira, Confluence, Regex, HTML, Splunk, Linux, MongoDB, Figma